

Daniel Gordon

gordondan2@gmail.com | +1 (647) 622-0424 | danielgordon.vercel.app | linkedin.com/in/daniel-gordon2 | github.com/DanielGordOn

SUMMARY

Software engineer targeting Forward Deployed and field engineering roles. I build and operate production systems end-to-end across cloud infrastructure, AI pipelines, and client deployments, working directly with stakeholders to scope, ship, and own outcomes in ambiguous environments.

EDUCATION

Wilfrid Laurier University Waterloo, ON
B.Sc. Honours Computer Science and Management

Sep 2022 – Dec 2026

CERTIFICATIONS

AWS Certified Cloud Practitioner Amazon Web Services
AWS Solutions Architect Associate Amazon Web Services

May 2026
In progress

EXPERIENCE

DevOps / Cloud Engineering Intern WellnessLiving Systems | Remote Apr 2026 – Present

- **Discovered a silent CloudWatch failure mode** where S3 request metrics returned empty data on 14 buckets due to disabled server-access logging; built a three-layer verification framework (GetRequests, PutRequests, Last Modified) that prevented 4 active production buckets from being incorrectly modified.
- **Owned end-to-end S3 cost optimization** across a 17-account AWS org, auditing 455 buckets (\$8,214/mo) via Python, executing an Intelligent-Tiering lifecycle rollout, and confirming 151,892 object transitions within 4 days with zero production incidents.
- **Hardened 79 production WordPress EC2 instances** fleet-wide via AWS SSM Run Command with zero direct server access, automating a 7-step security standard (Wordfence, PHP execution blocking, admin cleanup, Qualys) and delivering a 33-flag prioritized remediation report to stakeholders.
- **Executed a PHP 8.3 upgrade across 78 live client sites** via staged batches (10, 25, 42 instances); caught a missing php8.3-mysql extension after batch 1 before full-fleet impact, limiting client-facing downtime to under 5 minutes.
- **Authored a 14-section production SOP** for OS patch management across Amazon Linux 2023, Ubuntu 22.04, and Windows Server 2025, reviewed and approved by senior engineers.

Software Engineer, promoted to Lead MCG Digital Assetization | Remote Jul 2025 – May 2026

- **Promoted to technical lead after 6 months and integrated OpenAI Whisper and LLM orchestration** into production AI workflows for real-time 3D avatar animation, benchmarking accuracy and latency tradeoffs across models; improved overall system performance by 25%.

Research Assistant, Combinatorial Mathematics and HPC Wilfrid Laurier University | Waterloo, ON Sep 2025 – Present

- **Developed a high-performance C++ backtracking solver (BS45 Quantum Explorer)** deployed across 4 national HPC clusters (Trillium, Nibi, Fir, Rorqual) via SLURM, coordinating 10,000+ CPU cores to search for new Base Sequences, an unsolved open problem in combinatorial mathematics.
- **Implemented Theorem 2.3 residue pruning, spectral DFT filtering, and Wang-Zhu pair encoding** to reduce the naive $2^{(4n)}$ search space by ~100x; actively searching for BS(45,44), an unsolved world-record target.

Junior Solutions Architect LaunchPath Inc. | Toronto, ON Jun 2025 – Aug 2025

- **Designed a resilient dual-model Gemini AI gateway** using native fetch with AbortSignal timeouts and automatic failover between models; engineered structured prompt chains constraining the LLM to output strict JSON with step-by-step mathematical reasoning for deterministic outcomes.
- **Built a dynamic 15-step conditional questionnaire engine in React Native** where the steps array mutates in real time based on user input while keeping UI state stable, following Agile development practices; backed by a custom Node.js/Express REST API with JWT/bcrypt authentication and email verification built from scratch.

Software Engineer, AI Systems Self-Employed Apr 2024 – May 2026

- **Designed and shipped Phoenix Bot**, a cross-platform Electron desktop app (macOS, Windows, Linux) automating WhatsApp group chat summarization via Google Gemini; solved the core challenge of sustaining a persistent Puppeteer/headless-Chrome session inside an Electron main process, with OS keychain credential storage and local SQLite for session state.

PROJECTS

TellTours — AI Tour Guide App (Co-founder) telltours.com | Pre-release, beta waitlist live 2024 – Present

- **Built a provider-agnostic AI narration pipeline** routing between Gemini, OpenAI, and Ollama via a unified hot-swappable client, with geofenced playback triggers, in-flight LLM observability logging (model, token count, latency), and AI prompts externalized to Firestore for zero-redeploy behavior tuning.
- **Architected on GCP** (Cloud Run, Firestore, Cloud TTS) with two isolated environments and keyless CI/CD via GitHub Actions and Workload Identity Federation.

TECHNICAL SKILLS

Languages: Python, JavaScript / TypeScript, C++, Java, SQL, Bash, PowerShell

Cloud: AWS (S3, EC2, SSM, IAM, CloudWatch, Lambda, EBS, Identity Center), GCP (Cloud Run, Firestore, Compute Engine), Docker, GitHub Actions

DevOps: AWS SSM (Run Command, Patch Manager, Fleet Manager, State Manager), fleet automation, AWS CLI, incident response, FinOps, SLURM / HPC

Backend and AI: REST API design and integration, microservices, LLM integration (Gemini, Whisper, OpenAI), real-time pipelines, distributed systems, system design, React Native, Node.js, Electron