

Daniel Gordon

gordondan2@gmail.com | +1 (647) 622-0424 | danielgordon.vercel.app | linkedin.com/in/daniel-gordon2 | github.com/DanielGordOn

SUMMARY

Software engineer targeting Forward Deployed and field engineering roles. I build and operate production systems end-to-end across cloud infrastructure, AI pipelines, and client deployments, working directly with stakeholders to scope, ship, and own outcomes in ambiguous environments.

EDUCATION

Wilfrid Laurier University Waterloo, ON Sep 2022 – Dec 2026

B.Sc. Honours Computer Science and Management | Exploring double degree extension: B.Sc. CS + BBA

CERTIFICATIONS

AWS Certified Cloud Practitioner Amazon Web Services May 2026

AWS Solutions Architect Associate Amazon Web Services In progress

EXPERIENCE

DevOps / Cloud Engineering Intern WellnessLiving Systems | Remote Apr 2026 – Present

- **Discovered a silent CloudWatch failure mode** where S3 request metrics returned empty data on 14 buckets due to disabled server-access logging; built a three-layer verification framework (GetRequests, PutRequests, Last Modified) that prevented 4 active production buckets from being incorrectly modified.
- **Owned end-to-end S3 cost optimization** across a 17-account AWS org, auditing 455 buckets (\$8,214/mo) via Python, executing an Intelligent-Tiering lifecycle rollout, and confirming 151,892 object transitions within 4 days with zero production incidents.
- **Hardened 79 production WordPress EC2 instances** fleet-wide via AWS SSM Run Command with zero direct server access, automating a 7-step security standard (Wordfence, PHP execution blocking, admin cleanup, Qualys) and delivering a 33-flag prioritized remediation report to stakeholders.
- **Executed a PHP 8.3 upgrade across 78 live client sites** via staged batches (10, 25, 42 instances); caught a missing php8.3-mysql extension after batch 1 before full-fleet impact, limiting client-facing downtime to under 5 minutes.
- **Authored a 14-section production SOP** for OS patch management across Amazon Linux 2023, Ubuntu 22.04, and Windows Server 2025, reviewed and approved by senior engineers.

Software Engineer, promoted to Lead MCG Digital Assetization | Remote Jul 2025 – May 2026

- **Promoted to technical lead after 6 months and integrated OpenAI Whisper and LLM orchestration** into production AI workflows for real-time 3D avatar animation, benchmarking accuracy and latency tradeoffs across models; improved overall system performance by 25%.

Research Assistant, Combinatorial Mathematics and HPC Wilfrid Laurier University | Waterloo, ON Sep 2025 – Present

- **Developed a high-performance C++ backtracking solver (BS45 Quantum Explorer)** deployed across 4 national HPC clusters (Trillium, Nibi, Fir, Rorqual) via SLURM, coordinating 10,000+ CPU cores to search for new Base Sequences, an unsolved open problem in combinatorial mathematics.
- **Implemented Theorem 2.3 residue pruning, spectral DFT filtering, and Wang-Zhu pair encoding** to reduce the naive 2^{4n} search space by $\sim 100\times$; actively searching for BS(45,44), an unsolved world-record target.

Junior Solutions Architect LaunchPath Inc. | Toronto, ON Jun 2025 – Aug 2025

- **Designed a resilient dual-model Gemini AI gateway** using native fetch with AbortSignal timeouts and automatic failover between models; engineered structured prompt chains constraining the LLM to output strict JSON with step-by-step mathematical reasoning for deterministic outcomes.
- **Built a dynamic 15-step conditional questionnaire engine in React Native** where the steps array mutates in real time based on user input while keeping UI state stable; backed by a custom Node.js/Express API with JWT/bcrypt authentication and email verification built from scratch.

Software Engineer, AI Systems Self-Employed Apr 2024 – May 2026

- **Designed and shipped Phoenix Bot**, a cross-platform Electron desktop app (macOS, Windows, Linux) automating WhatsApp group chat summarization via Google Gemini; solved the core challenge of sustaining a persistent Puppeteer/headless-Chrome session inside an Electron main process, with OS keychain credential storage and local SQLite for session state.

PROJECTS

TellTours — AI Tour Guide App (Co-founder) telltours.com | Pre-release, beta waitlist live 2024 – Present

- **Built a provider-agnostic AI narration pipeline** routing between Gemini, OpenAI, and Ollama via a unified hot-swappable client, with geofenced playback triggers, in-flight LLM observability logging (model, token count, latency), and AI prompts externalized to Firestore for zero-redeploy behavior tuning.
- **Architected on GCP** (Cloud Run, Firestore, Cloud TTS) with two isolated environments and keyless CI/CD via GitHub Actions and Workload Identity Federation.

TECHNICAL SKILLS

Languages: Python, JavaScript / TypeScript, C++, Java, SQL, Bash, PowerShell

Cloud: AWS (S3, EC2, SSM, IAM, CloudWatch, Lambda, EBS, Identity Center), GCP (Cloud Run, Firestore, Compute Engine), Docker, GitHub Actions

DevOps: AWS SSM (Run Command, Patch Manager, Fleet Manager, State Manager), fleet automation, AWS CLI, incident response, FinOps, SLURM / HPC

Backend and AI: REST APIs, microservices, LLM integration (Gemini, Whisper, OpenAI), real-time pipelines, distributed systems, system design, React Native, Node.js, Electron